

ALGORITHMS – COMTEK3, CCT3 & ESD3


ORDER STATISTIC, BINARY SEARCH TREES AND DATA AUGMENTATION

Ramoni Adeogun

Associate Professor & Head AI for Communications
Wireless Communication Networks Section (WCN)
Department of Electronic Systems
Email: ra@es.aau.dk

Agenda

- **Order statistics**
 - Find maximum and minimum
 - Find median
- **Binary search trees**
 - Structure of trees
 - Operations to add, remove and search
- **Red-black trees**
 - Structure of trees
 - Operations to add, remove and search
 - Property maintenance
- **Data structure augmentation**
- Exercises



Anatomy of a binary search tree

Image source: <https://algs4.cs.princeton.edu/32bst/>

Order Statistics

- The i th **order statistic** of a set of n elements is the i th smallest element.
- Find the i th smallest (i.e. of rank i) of n elements, $A[1], A[2], \dots, A[n]$, mutually comparable.
 - $i = 1$: the minimum, $i = n$: the maximum
 - Find both minimum and maximum
 - $i = (n + 1)/2$ is the unique median when n is odd, otherwise $i = \lfloor (n + 1)/2 \rfloor$ is the lower median and $i = \lceil (n + 1)/2 \rceil$ the upper median
- A simple solution is to sort and pick element i , but then we are doing unnecessary work.
- Count comparisons between elements and assume the elements are distinct, for simplicity.



Find maximum

- Initialize *max* to $A[1]$ and compare it to $A[2]$. If $A[2]$ is larger it becomes the new *max*.
- The procedure of comparing *max* to a new element and possibly updating *max* is repeated until all n elements have been examined.
- This takes $n - 1$ comparisons between elements, which cannot be reduced, since each element but one must be compared to *max* or to some element that is smaller than *max*.



Find maximum and minimum

- To find both the largest and the smallest of n elements:
 - we can first find the largest in $n - 1$ comparisons, and then the smallest among the remaining $n - 1$ elements in $n - 2$ comparisons;
 - leading to a total of $2n - 3$ comparisons in total.
 - Can we do any faster? Oh yes!!
- Solve the problem for n elements by
 - removing two arbitrary elements, x and y , and solving the problem recursively for the remaining $n - 2$ elements.
 - The algorithm will return the largest and smallest elements among these $n - 2$ elements.
 - To get the maximum and minimum:
 - compare the larger of x and y to largest of $n - 2$ and the smaller to the minimum of $n - 2$ elements.
- The number of comparisons needed is hence expressed by the recurrence:

$$T(n) = \begin{cases} 0 & \text{if } n = 1 \\ 1 & \text{if } n = 2 \\ T(n - 2) + 3 & \text{if } n > 2 \end{cases}$$

$T(n) = T(n - 4) + 3 + 3 = T(n - 6) + 3 + 3 + 3 \approx 3n/2$, or more precisely:

$$T(n - 2) + 3 = \begin{cases} 3(n - 2)/2 + T(2) = 3n/2 - 2 & \text{for even } n \\ 3(n - 1)/2 + T(1) = (3n + 1)/2 - 2 & \text{for odd } n \end{cases}$$

Selection in $O(n)$

- More difficult problem than finding a minimum:
 - e.g., find an element with a rank around $n/2$ in an n –element array
 - This can be used to partition a file in smaller *quantiles* (when only a small part of the data is needed for later processing).
 - We will study an algorithm to do this in linear time –
- **RANDOMIZED-SELECT:**
 - similar to Quicksort and has $O(n)$ average case time
 - calls RANDOMIZED-PARTITION:
 - randomly chooses a pivot element $A[q]$ and splits $A[p, \dots, r]$ so that elements less than the pivot are placed in $A[p, \dots, q - 1]$,
 - and elements larger are placed in $A[q + 1, \dots, r]$.
 - If the rank k (relative to p) of the pivot is greater than i ,
 - the search continues among the smaller elements,
 - otherwise among the larger elements.



```
RANDOMIZED-SELECT( $A, p, r, i$ )
1  if  $p == r$ 
2      return  $A[p]$       //  $1 \leq i \leq r - p + 1$  when  $p == r$  means that  $i = 1$ 
3   $q = \text{RANDOMIZED-PARTITION}(A, p, r)$ 
4   $k = q - p + 1$ 
5  if  $i == k$ 
6      return  $A[q]$       // the pivot value is the answer
7  elseif  $i < k$ 
8      return RANDOMIZED-SELECT( $A, p, q - 1, i$ )
9  else return RANDOMIZED-SELECT( $A, q + 1, r, i - k$ )
```



Binary Search Trees

Binary Search Trees

- **Binary search trees** are an important data structure for dynamic sets.
 - Accomplish many dynamic-set operations in $O(h)$ time, where h = height of the tree
 - We represent a binary tree by a linked data structure in which each node is an object.
 - T.root points to the root of tree T
 - Each node contains the attributes
 - key (and possibly other satellite data).
 - left: points to left child.
 - right: points to right child.
 - p: points to parent. T.root.p = NIL
- Stored keys must satisfy the binary-search-tree property.
 - If y is a node in the left subtree of x , then $y.key \leq x.key$
 - If y is a node in the right subtree of x , then $y.key > x.key$



INORDER-TREE-WALK

- The binary-search-tree property allows us to print keys in a binary search tree in order, recursively, using an algorithm called an **INORDER-TREE-WALK**.
 - Elements are printed in monotonically increasing order
- How INORDER-TREE-WALK works:
 - Check to make sure that x is not NIL.
 - Recursively, print the keys of the nodes in x 's left subtree.
 - Print x 's key.
 - Recursively, print the keys of the nodes in x 's right subtree.
- **Example:** Apply INORDER-TREE-WALK to the trees on the previous slide
- **Time:**
 - Intuitively, the walk takes $\Theta(n)$. Why?

```
INORDER-TREE-WALK( $x$ )  
1  if  $x \neq \text{NIL}$   
2      INORDER-TREE-WALK( $x.\text{left}$ )  
3      print  $x.\text{key}$   
4      INORDER-TREE-WALK( $x.\text{right}$ )
```



Operations in binary search trees

- Types of operations
 - Looking for a specific key
 - Minimum
 - Maximum
 - Next
 - Previous
- All can be carried out in $O(h)$ time with h as the height of the tree.
- In worst case, the tree is just n elements high, and in best $\log_2(n)$



Search in a binary tree

Goal: Search for a node in a BST

Inputs:

x – a pointer to the root

k – key for the node to search

Examples:

$k = 13$: $15 \rightarrow 6 \rightarrow 7 \rightarrow 13$


$k = 4$: $15 \rightarrow 6 \rightarrow 3 \rightarrow 4$

TREE-SEARCH(x, k)

```
1  if  $x == \text{NIL}$  or  $k == x.\text{key}$ 
2      return  $x$ 
3  if  $k < x.\text{key}$ 
4      return TREE-SEARCH( $x.\text{left}, k$ )
5  else return TREE-SEARCH( $x.\text{right}, k$ )
```

ITERATIVE-TREE-SEARCH(x, k)

```
1  while  $x \neq \text{NIL}$  and  $k \neq x.\text{key}$ 
2      if  $k < x.\text{key}$ 
3           $x = x.\text{left}$ 
4      else  $x = x.\text{right}$ 
5  return  $x$ 
```



Minimum and maximum


- The binary-search-tree property guarantees that
 - the minimum key of a binary search tree is located at the leftmost node, and
 - the maximum key of a binary search tree is located at the rightmost node.
- Traverse the appropriate pointers (left or right) until NIL is reached.

TREE-MINIMUM(x)

```
1 while  $x.left \neq \text{NIL}$ 
2    $x = x.left$ 
3 return  $x$ 
```

TREE-MAXIMUM(x)

```
1 while  $x.right \neq \text{NIL}$ 
2    $x = x.right$ 
3 return  $x$ 
```



- Both procedures visit nodes that form a downward path from the root to a leaf.
 - Both procedures run in $O(h)$ time, where h is the height of the tree.



Successor and predecessor



Assuming that all keys are distinct, the successor of a node x is the node y such that $y.key$ is the smallest key $> x.key$. If x has the largest key in BST, then we say that x 's successor is NIL.

There are two cases:

1. If node x has a non-empty right subtree, then x 's successor is the minimum in x 's right subtree.
2. If node x has an empty right subtree, notice that:
 - If we move to the left up the tree (move up through right children), we're visiting smaller keys.
 - x 's successor y is the node that x is the predecessor of (x is the maximum in y 's left subtree).

TREE-SUCCESSOR(x)

```
1 if  $x.right \neq \text{NIL}$ 
2   return TREE-MINIMUM( $x.right$ ) //
3 else // find the lowest ancestor of  $x$  whose
4    $y = x.p$ 
5   while  $y \neq \text{NIL}$  and  $x == y.right$ 
6      $x = y$ 
7      $y = y.p$ 
8   return  $y$ 
```



- Note: TREE-PREDECESSOR is symmetric to TREE-SUCCESSOR




Insert into a BST

- Simple to insert, but we must retain the property of the binary tree
- Go through the tree and find an empty spot that matches criteria
 - Keys must obey binary tree property
- Update pointers

TREE-INSERT(T, z)

```
1   $x = T.root$            // node being compared with  $z$ 
2   $y = NIL$               //  $y$  will be parent of  $z$ 
3  while  $x \neq NIL$       // descend until reaching a leaf
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$               // found the location—insert  $z$  with parent  $y$ 
9  if  $y == NIL$ 
10      $T.root = z$        // tree  $T$  was empty
11 elseif  $z.key < y.key$ 
12      $y.left = z$ 
13 else  $y.right = z$ 
```



Example: Insert node z with $key[z]=13$



Deleting nodes from a BST

- Conceptually, deleting node z from binary search tree T has three cases:
 1. If z has no children, just remove it
 2. If z has just one child, then make that child take z 's position in the tree, dragging the child's subtree along.
 3. If z has two children, then find z 's successor y and replace z by y in the tree.
 - y must be in z 's right subtree and have no left child.
 - The rest of z 's original right subtree becomes y 's new right subtree, and z 's left subtree becomes y 's new left subtree.

TRANSPLANT(T, u, v)

```
1  if  $u.p == \text{NIL}$ 
2       $T.root = v$ 
3  elseif  $u == u.p.left$ 
4       $u.p.left = v$ 
5  else  $u.p.right = v$ 
6  if  $v \neq \text{NIL}$ 
7       $v.p = u.p$ 
```





TREE-DELETE(T, z)

```
1  if  $z.left == \text{NIL}$ 
2      TRANSPLANT( $T, z, z.right$ )
3  elseif  $z.right == \text{NIL}$ 
4      TRANSPLANT( $T, z, z.left$ )
5  else  $y = \text{TREE-MINIMUM}(z.right)$ 
6      if  $y \neq z.right$ 
7          TRANSPLANT( $T, y, y.right$ )
8           $y.right = z.right$ 
9           $y.right.p = y$ 
10     TRANSPLANT( $T, z, y$ )
11      $y.left = z.left$ 
12      $y.left.p = y$ 
```



Deleting from a BST (2)

1. If z has no left child, replace z by its right child. The right child may or may not be NIL.
2. If z has just one child, and that child is its left child, then replace z by its left child.
3. Otherwise, z has two children. Find z 's successor y . y must lie in z 's right subtree and have no left child
 - I. If y is z 's right child, replace z by y and leave y 's right child alone.
 - II. Otherwise, y lies within z 's right subtree but is not the root of this subtree. Replace y by its own right child. Then replace z by y .



Red Black Trees



Red-black Trees

- A red-black tree is a binary search tree with an extra attribute: an attribute color, which is either red or black.
- All leaves are empty (nil) and colored black.
 - We use a single sentinel, $T.nil$, for all the leaves of red-black tree T .
 - $T.nil.color$ is black.
 - The root's parent is also $T.nil$.
- All other attributes of binary search trees are inherited by red-black trees (key, left, right, and p). We don't care about the key in $T.nil$.



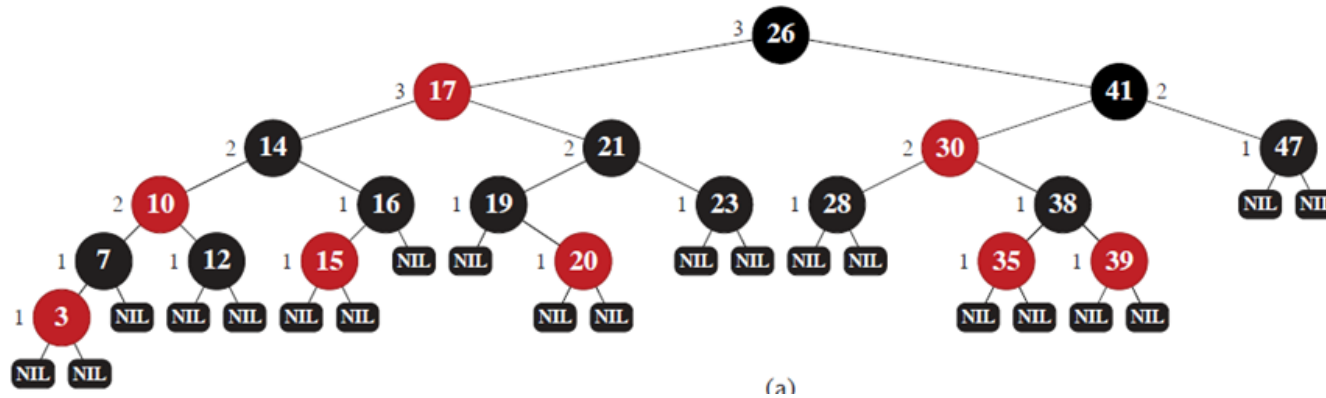
Red-black tree properties

1. Every node is **either red or black**.
2. The **root is black**.
3. Every **leaf (T.nil) is black**.
4. If a node is **red**, then **both its children are black** => no two reds in a row on a simple path from the root to a leaf.
5. For each node, **all paths from the node to descendant leaves contain the same number of black nodes**.



Height of a red-black tree

- **Height of a node** is the number of edges in a longest path to a leaf.
- **Black-height of a node x :** $bh(x)$ is the number of black nodes (including T.nil) on the path from x to leaf, not counting x .
 - By property 5, black-height is well defined.



- **Claim:** Any node with height h has black-height $\geq h/2$
 - Proof: By property 4, $\leq h/2$ nodes on the path from the node to a leaf are red. Hence $\geq h/2$ nodes are black




Operations on red-black trees

- The non-modifying binary-search-tree operations MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR, and SEARCH run in $O(\text{height})$ time. Thus, they take $O(\log_2(n))$ time on red-black trees.
- Insertion and deletion are not so easy.
 - If we insert, what color to make the new node?
 - Red? Might violate property 4.
 - Black? Might violate property 5.
 - If we delete, thus removing a node, what color was the node that was removed?
 - Red? OK, since we won't have changed any black-heights, nor will we have created two red nodes in a row. Also, cannot cause a violation of property 2, since if the removed node was red, it could not have been the root.
 - Black? Could cause there to be two reds in a row (violating property 4), and can also cause a violation of property 5. Could also cause a violation of property 2, if the removed node was the root and its child—which becomes the new root—was red.



Rotations

- The basic tree-restructuring operation needed to maintain red-black trees as balanced binary search trees.
- Changes the local pointer structure. (Only pointers are changed.)
- Have both left rotation and right rotation. They are inverses of each other.
- A rotation takes a red-black-tree and a node within the tree.



- The pseudocode for LEFT-ROTATE assumes that
 - $x.right \neq T.nil$, and
 - root's parent is $T.nil$.
- Pseudocode for RIGHT-ROTATE is symmetric: exchange left and right everywhere

```
LEFT-ROTATE( $T, x$ )
1   $y = x.right$ 
2   $x.right = y.left$ 
3  if  $y.left \neq T.nil$ 
4       $y.left.p = x$ 
5   $y.p = x.p$ 
6  if  $x.p == T.nil$ 
7       $T.root = y$ 
8  elseif  $x == x.p.left$ 
9       $x.p.left = y$ 
10 else  $x.p.right = y$ 
11  $y.left = x$ 
12  $x.p = y$ 
```




Rotations: Example

- **Before rotation:** keys of x 's left subtree $\leq 11 \leq$ keys of y 's left subtree $\leq 18 \leq$ keys of y 's right subtree.
- Rotation makes y 's left subtree into x 's right subtree.
- **After rotation:** keys of x 's left subtree $\leq 11 \leq$ keys of x 's right subtree $\leq 18 \leq$ keys of y 's right subtree.

Time: $O(1)$ for both LEFT-ROTATE & RIGHT-ROTATE

- a constant number of pointers are modified.



Insertion

- Similar to insertion in regular binary-search-tree insertion.
 - RB-INSERT ends by coloring the new node z red.
 - Then it calls RB-INSERT-FIXUP because we could have violated a red-black property.
 - Which property might be violated
 1. OK
 2. If z is the root, then there's a violation. Otherwise, OK.
 3. OK.
 4. If $z.p$ is red, there's a violation: both z and $z.p$ are red.
 5. OK.
 - Remove the violation by calling RB-INSERT-FIXUP.
- ```
RB-INSERT(T, z)
1 $x = T.root$ // node being compared with z
2 $y = T.nil$ // y will be parent of z
3 while $x \neq T.nil$ // descend until reaching the sentinel
4 $y = x$
5 if $z.key < x.key$
6 $x = x.left$
7 else $x = x.right$
8 $z.p = y$ // found the location—insert z with parent y
9 if $y == T.nil$
10 $T.root = z$ // tree T was empty
11 elseif $z.key < y.key$
12 $y.left = z$
13 else $y.right = z$
14 $z.left = T.nil$ // both of z 's children are the sentinel
15 $z.right = T.nil$
16 $z.color = RED$ // the new node starts out red
17 RB-INSERT-FIXUP(T, z) // correct any violations of red-black properties
```





# RB-Insert Fix-up: Example

Loop invariant:

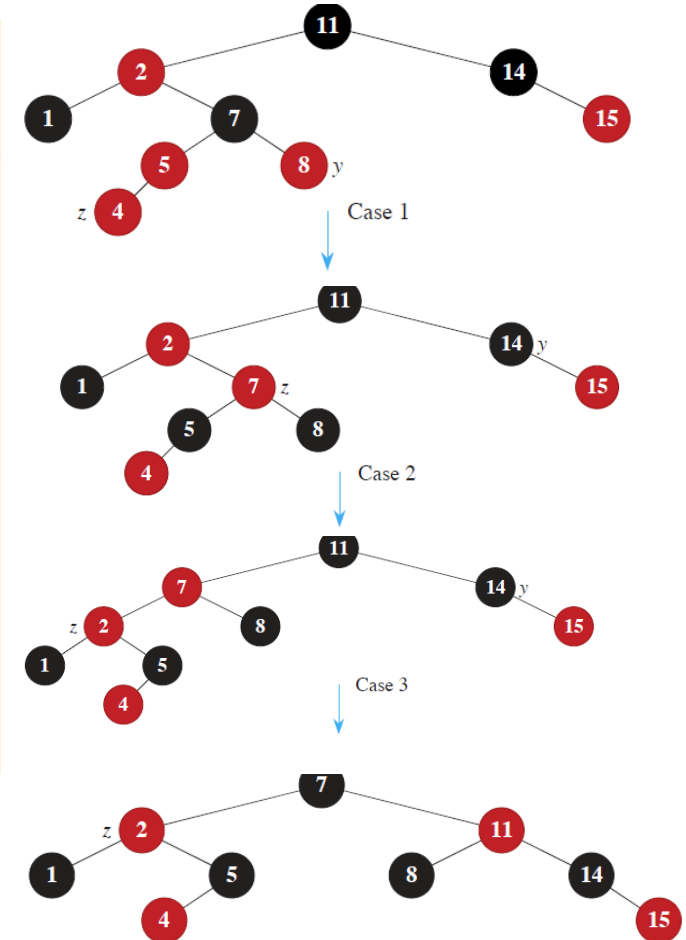
- Node  $z$  is red
- If  $z.p$  is the root, then  $z.p$  is black.
- There is at most one red-black violation:
  - Property 2:  $z$  is a red root, or
  - Property 4:  $z$  and  $z.p$  are both red.
- Analysis:
- $O(\log_2(n))$  time to get through RB-INSERT up to the call of RB-INSERT-FIXUP.

RB-INSERT-FIXUP( $T, z$ )

```

1 while $z.p.color == RED$
2 if $z.p == z.p.p.left$ // is z 's parent a left child?
3 $y = z.p.p.right$ // y is z 's uncle
4 if $y.color == RED$ // are z 's parent and uncle both red?
5 $z.p.color = BLACK$
6 $y.color = BLACK$
7 $z.p.p.color = RED$
8 $z = z.p.p$
9 } case 1
10 else
11 if $z == z.p.right$
12 $z = z.p$
13 LEFT-ROTATE(T, z)
14 $z.p.color = BLACK$
15 $z.p.p.color = RED$
16 RIGHT-ROTATE($T, z.p.p$)
17 } case 2
18 else // same as lines 3–15, but with "right" and "left" exchanged
19 $y = z.p.p.left$
20 if $y.color == RED$
21 $z.p.color = BLACK$
22 $y.color = BLACK$
23 $z.p.p.color = RED$
24 $z = z.p.p$
25 } case 3
26 else
27 if $z == z.p.left$
28 $z = z.p$
29 RIGHT-ROTATE(T, z)
30 $z.p.color = BLACK$
31 $z.p.p.color = RED$
32 LEFT-ROTATE($T, z.p.p$)
33 $T.root.color = BLACK$

```



# Deletion in RB-Trees

- Based on the TREE-DELETE procedure for BSTs:

```

RB-DELETE-FIXUP(T, x)
1 while x ≠ T.root and x.color == BLACK
2 if x == x.p.left // is x a left child?
3 w = x.p.right // w is x's sibling
4 if w.color == RED
5 w.color = BLACK
6 x.p.color = RED
7 LEFT-ROTATE(T, x.p)
8 w = x.p.right
9 if w.left.color == BLACK and w.right.color == BLACK
10 w.color = RED
11 x = x.p
12 else
13 if w.right.color == BLACK
14 w.left.color = BLACK
15 w.color = RED
16 RIGHT-ROTATE(T, w)
17 w = x.p.right
18 w.color = x.p.color
19 x.p.color = BLACK
20 w.right.color = BLACK
21 LEFT-ROTATE(T, x.p)
22 x = T.root
23 else // same as lines 3–22, but with "right" and "left" exchanged
24 w = x.p.left
25 if w.color == RED
26 w.color = BLACK
27 x.p.color = RED
28 RIGHT-ROTATE(T, x.p)
29 w = x.p.left
30 if w.right.color == BLACK and w.left.color == BLACK
31 w.color = RED
32 x = x.p
33 else
34 if w.left.color == BLACK
35 w.right.color = BLACK
36 w.color = RED
37 LEFT-ROTATE(T, w)
38 w = x.p.left
39 w.color = x.p.color
40 x.p.color = BLACK
41 w.left.color = BLACK
42 RIGHT-ROTATE(T, x.p)
43 x = T.root
44 x.color = BLACK

```

RB-TRANSPLANT(T, u, v)

```

1 if u.p == T.nil
2 T.root = v
3 elseif u == u.p.left
4 u.p.left = v
5 else u.p.right = v
6 v.p = u.p

```

RB-DELETE(T, z)

```

1 y = z
2 y-original-color = y.color
3 if z.left == T.nil
4 x = z.right
5 RB-TRANSPLANT(T, z, z.right) // replace z by its right child
6 elseif z.right == T.nil
7 x = z.left
8 RB-TRANSPLANT(T, z, z.left) // replace z by its left child
9 else y = TREE-MINIMUM(z.right) // y is z's successor
10 y-original-color = y.color
11 x = y.right
12 if y ≠ z.right // is y farther down the tree?
13 RB-TRANSPLANT(T, y, y.right) // replace y by its right child
14 y.right = z.right // z's right child becomes
15 y.right.p = y // y's right child
16 else x.p = y // in case x is T.nil
17 RB-TRANSPLANT(T, z, y) // replace z by its successor y
18 y.left = z.left // and give z's left child to y,
19 y.left.p = y // which had no left child
20 y.color = z.color
21 if y-original-color == BLACK // if any red-black violations occurred,
22 RB-DELETE-FIXUP(T, x) // correct them

```



# Dynamic Order Statistics and Data Augmentation


# Augmenting Data Structures

- Suppose we have a base data structure  $D$  that efficiently handles a standard set of operations. **e.g.,  $D$  is a RBT that supports operations SEARCH, INSERT, and DELETE**
- In some applications, we desire a data structure that support an additional set of operations. For example, **the order-statistics operations SELECT and RANK**
- So, how do we efficiently implement the new operations without degrading the existing ones?
- By augmenting the data structure:
  - Choose underlying data structure, for instance a red-black tree.
  - Determine additional information to be maintained, for instance sizes of subtrees.
  - Verify that additional information is updated correctly for the operations on the data structure.
  - Develop new operations.



# Dynamic Order Statistics

- Require other operations in addition to standard dynamic set operations
  - OS-SELECT( $x, i$ ) – returns  $i$ th smallest key in subtree rooted at  $x$
  - OS-RANK( $T, x$ ) – returns rank of  $x$  in the linear order determined by an inorder traversal of  $T$
- *Idea:* Store sizes of subtrees in the nodes in a red-black tree  $\Rightarrow$  **order-statistic tree** .



# Operations

- OS-SELECT for retrieving an element with a given rank
- OS-RANK for determining the rank of an element
- Both OS-SELECT and OS-RANK takes  $O(\log n)$  time.

OS-SELECT( $x, i$ )

```
1 $r = x.left.size + 1$ // rank of x within the subtree rooted at x
2 if $i == r$
3 return x
4 elseif $i < r$
5 return OS-SELECT($x.left, i$)
6 else return OS-SELECT($x.right, i - r$)
```

OS-RANK( $T, x$ )

```
1 $r = x.left.size + 1$ // rank of x within the subtree rooted at x
2 $y = x$ // root of subtree being examined
3 while $y \neq T.root$
4 if $y == y.p.right$ // if root of a right subtree ...
5 $r = r + y.p.left.size + 1$ // ... add in parent and its left subtree
6 $y = y.p$ // move y toward the root
7 return r
```



# Summary

- **Binary search trees** are an important data structure for dynamic sets.
  - Accomplish many dynamic-set operations in  $O(h)$  time, where  $h$  = height of tree.
  - $T.root$  points to the root of tree  $T$ .
  - Each node contains the attributes: key (and possibly other satellite data), left - points to left child, right - points to right child and  $p$  - points to parent.  $T.root.p = NIL$
- **Red black trees:**
  - A variation of binary search trees with colours – red or black attribute.
  - Balanced: height is  $O(\log_2(n))$  where  $n$  is the number of nodes.
  - Operations will take  $O(\log_2(n))$  time in the worst case.
- **Order Statistic:**
  - $i$ th order statistic is the  $i$ th smallest element of a set of  $n$  elements
  - Applications in finding minimum, maximum, median, selection, etc
- **Data augmentation:**
  - Principle used in applications in which we desire a data structure that support an additional set of operations beyond a base structure

